

# Fraud Detection Analysis

## Business Context

This notebook presents a solution for predicting fraudulent transactions. The goal is to develop a machine learning model and provide actionable insights.

## Candidate Expectations & Answers

This notebook is structured to address the following:

1. Data Cleaning & EDA
2. Model Description
3. Variable Selection
4. Model Performance
5. Key Predictive Factors
6. Factor Interpretation
7. Prevention Strategies
8. Action Evaluation

### 1. Data Loading and Imports

### 2. Data Cleaning and EDA

#### **Question 1: Data cleaning including missing values, outliers and multi-collinearity.**

We investigate missing values, check for outliers in numerical columns, and analyze multi-collinearity.

#### **Observations:**

- No missing values found in the dataset.
- **Outliers:** The `amount` column has a massive range (max ~92M vs mean ~180k). Similarly for balances. These are likely valid large transactions rather than errors, but they are outliers in the statistical sense.
- **Multi-collinearity:** We expect `oldbalanceOrg` and `newbalanceOrig` to be highly correlated.

### 3. Feature Engineering and Variable Selection

### **Question 3: How did you select variables to be included in the model?**

#### **Selection Strategy:**

1. **Included:** `step`, `type`, `amount`, `oldbalanceOrg`, `newbalanceOrig`, `oldbalanceDest`, `newbalanceDest`. These provide the core transaction details.
2. **Excluded:** `nameOrig`, `nameDest`. High cardinality categorical variables. While specific accounts might be repeat offenders, for a generalizable model, we initially exclude them to avoid overfitting to specific IDs. `isFlaggedFraud` is also excluded to prevent leakage if it's a post-event flag, or it can be used as a baseline comparison.
3. **Feature Engineering:**
  - `type` : Converted to numerical using Label Encoding.
  - `errorBalanceOrig` : Difference between original balance difference and transaction amount.
  - `errorBalanceDest` : Difference between destination balance difference and transaction amount.

## **4. Model Development**

### **Question 2: Describe your fraud detection model in elaboration.**

We use **XGBoost (Extreme Gradient Boosting)**. Reasons:

- Handling Large Data: Efficient matrix operations.
- Non-Linearity: Tree-based models capture complex interactions between balance and amount.
- Class Imbalance: XGBoost has built-in `scale_pos_weight` to handle imbalance.
- Interpretability: Provides feature importance scores.

We split the data 80% Training (Calibration) and 20% Testing (Validation).

## **5. Model Evaluation**

### **Question 4: Demonstrate the performance of the model by using best set of tools.**

We evaluate using Precision, Recall, F1-Score, and AUC-ROC.

## **6. Interpretability and Insights**

### **Question 5: What are the key factors that predict fraudulent customer? Question 6: Do these factors make sense? If yes, How? If not, How not?**

**Analysis of Factors:** The model typically highlights:

1. **errorBalanceOrig/Dest**: Large discrepancies in expected balance updates are strong indicators.
2. **amount**: Fraudulent transactions are often large efforts to cash out.
3. **type**: Transfer and Cash-Out are the primary vectors for fraud.

**Do they make sense? Yes.** Fraudsters aim to steal money (high amount), often move it (Transfer) and withdraw it (Cash-Out). They might leave the account in an inconsistent state or drain it completely (`oldbalanceOrg = amount`), leading to specific balance patterns.

## 7. Actionable Plan

**Question 7: What kind of prevention should be adopted while company update its infrastructure?**

1. **Real-time Transaction Scoring**: Deploy this XGBoost model to score transactions in real-time. Block or flag transactions with score > threshold (e.g., 0.9).
2. **Velocity Checks**: Limit the number/amount of transactions within a time window (using the `step` variable intuition).
3. **Two-Factor Authentication (2FA)**: Trigger 2FA for high-risk transactions (e.g., Transfers > \$10,000).
4. **Balance Validation**: Implement strict ACID compliance and checks to ensure `newbalance = oldbalance - amount`. Any deviation should auto-freeze.

**Question 8: Assuming these actions have been implemented, how would you determine if they work?**

1. **A/B Testing**: Run the model in "shadow mode" first, then on a subset of users.
2. **Metric Monitoring**:
  - **Fraud Rate**: Should decrease.
  - **False Positive Rate (Customer Friction)**: Should remain low. If too many legitimate users are blocked, adjust the threshold.
  - **Chargeback Rates**: A lagging indicator that should drop over months.